

Database Systems (8/22/2023)

Database Requirements

- Collection
- Universe of discourse
- Logically coherent
- Specific purpose
- Source(s)
- Interaction with events
- Actively interested audience
- Generated and maintained manually or by machine
- Any size and complexity

DBMS (Database Management System)

- Collection of general purpose software programs
- 2 functions (define the database, manipulate data)

	Simple and complex data	Simple data	Complex data
Query	(SQL)	(CQL) Object	- atomic
Relational	relational	- no manipulation	- resulting from calc
Object	(SQL!)	(CQL) Object	- computed
File System	Oriented	Ex: number, word	

Queried systems

- returns dynamic results

Ex: list of records for data range

Not-Queried Systems

- provides static view of data records

Ex: Text files that aren't searchable.

When not to build database

- costs increase for same result
- time to execute procedures increases
- lacking control & restriction
- difficult concepts & understanding

Database Systems (Ch. 1)

miniworld universe of discourse - some aspect of real world

meta data - DBMS in form of catalog or dictionary

query - causes some data to be retrieved

transaction - data read/written into database

- * data records are storing same type, data elements are the categories, data type is string, int, etc

requirements specification/analysis → conceptual design

physical design ← logical design

NOSQL systems - contain self describing data with names!

values in one structure

operation (function/method) contains interface & method

- * users have different view of database; a subset of database or containing virtual data not explicitly stored

concurrency control - ensures the result of several users updating same data is updated correctly

✓ online transaction processing (OLTP)

transaction - process that includes one or more database accesses

isolation - allows each transaction to execute in isolation of other transactions

* administrators, designers, end users, analysts, programmers

redundancy controlling - same data several times wastes space & is inconsistent

- * storing in one place is data normalization, putting a bunch of data into one file is denormalization

- * most DBMS do their own data buffering due to

Database Systems (8/24/2025)

3 categories of Data Models

- High level / conceptual model
- Low level / physical model
- Representational / implementation model

Schema construct - formal notation to describe entity w/ attributes

ENTITY = OBJECT = TABLE Attribute = Property = Column = Field

INSTANCE = Constructs + Domain = STATE

* changes when add, delete, modify

Database Systems (Ch. 2)

Entity - real world object/concept described through
attributes

access path - search structure for efficiency

description of database is database schema, displayed schema
is schema diagram

object in schema is schema construct

data definition language (DDL) & storage definition language (SDL)

3 Schema architecture & data manipulation language (DML)

* interfaces differ when provided by DBMS

database & DBMS catalog are stored on disk, accessed by OS,
typically with buffer management

* page 43 (PDF 74)

database utilities: loading, backup, database

storage recognition, performance monitoring

* accessing from remote location, connection
communication networks

* two tier vs three tier architectures

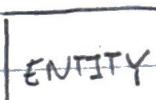
- three tier adds an intermediate layer
between client and database server (application or web server)

* classification of DBMS

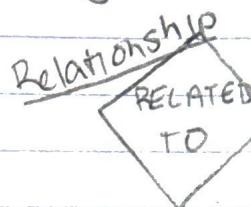
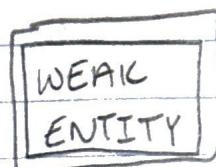
Database Systems (8/29/2023)

ER diagrams are a formal way of representing:

- objects
- properties
- relationships
- constraints

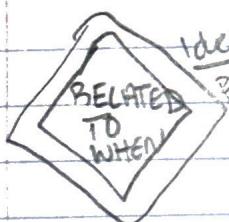


- rectangle with a name (usually a noun)

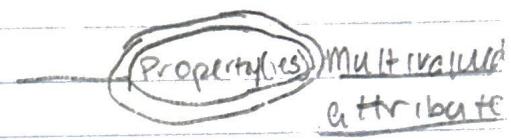


- connects two entities (often a verb)

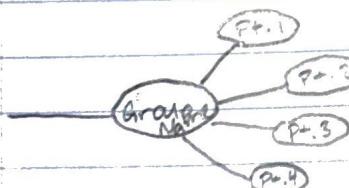
* never enter a baby without a mother



Atomic attribute



(P(n)) → Derived - complex
attribute data.



Composite - general atomic parts must be collected together

The Role — connects entity with relationship

The mandatory role = MUST participate in relationship

(min, max) constraint (min, max) - tells what the min & max participation in a relationship might be

The cardinality constraint (C) - shows # of members in entity + that can be related to members in another entity

Database Systems (8/29/2023)

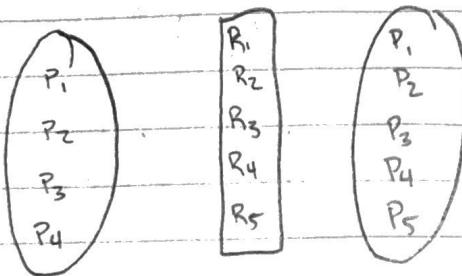
key

attribute - property that is always unique in value

F=Key

attribute - not necessarily unique, but used as matching value

* order of steps for constructing diagram



Database Systems (Ch. 3)

entity - thing or object

attributes - particular properties they describe

* EMPLOYEE entity might describe name, age, address, salary, job

simple (atomic) attributes - attributes no longer divisible

single-valued - attributes have single value for entity

derived attribute - derivable from stored attribute

(you can get age from birthdate)

* complex attributes

* relationships (binary & ternary, recursive relationships, cardinality ratio)

* UML class diagrams

* ER model on page 85, 89, 91, 94, 98, 99, 102

Database Systems (Ch. 16)

- * Primary storage, secondary storage, tertiary storage
 - Cache memory, DRAM, magnetic disks, CD-ROM, DVD, tapes
 - Flash memory, optical drives, magnetic tapes
- * secondary storage devices are nonvolatile, main memory is volatile
 - Data access more efficient on disk: buffering, organization, reading ahead of request, scheduling I/O, log disks, recovery purpose
 - pointer and then spanned
- * accessing disk block is expensive because of seek time, rotational delay, block transfer time
- * double buffering when accessing continuous disk blocks
- * fixed/variable length, spanned/unspanned, file header
- * 3 file organizations: ordered, unordered, hashed

Database Systems (9/5/2023)

Disks - works with read/write head to store/retrieve data

Tracks - concentric circles on disk

Cylinder - tracks with same diameter on consecutive disks

Blocks - range from 512 bytes to 4096 bytes; equal division of a track set by the operating system

Solid State Drives

- use electronic circuits to store/retrieve data; store data in blocks; no mechanical components; embedded processor
- SSDs read & write data to an underlying set of interconnect flash memory chips; store data even when it's not connected to a power source
- SSD's are able to read blocks of data that are spread out

Seek time, Block transfer time, Latency, Bulk Transfer Rate

$$\text{* Total Time} = \text{Seek} + \text{Latency} + \text{Block transfer time}$$

Large ~ 10 msec Large ~ 10 msec Small ~ 1 msec

Variable Length record - unknown size vs Fixed length

Mixed file - variable length field, multivalued attribute, optional

$$bFr = \text{floor}(B/R) \quad B = \text{block size in bytes} \quad R = \text{record size in bytes}$$

$$\text{UNUSED} = B - (bFr * R)$$

* red - record at a time blue - set at a time

Database Systems (9/17/2023)

- ① Look at the parent
- ② Determine if it has an attribute
 - a) Yes? What domain needs to be true.
 - b) No? Add it, indicate domain
- ③ make subclasses
- ④ Connect role 1:n to circle.

...

* 2 reasons to make subclasses (collect special properties, participate in a special relationship)

Database Systems (Ch.4)

Superclass/Subclass or supertype/subtype or class/subclass
Type inheritance - subclass inherits from class

Generalization - making a generalized entity from what's given (seeing truck and car then making vehicle class)
Predicate defined, Attribute defined, User defined

Disjointness constraint - entity can be a member of 1 subclass()

Overlapping constraint - entity can be member of >1 subclass(s)

Partial specialization - entity belongs to no subclasses

Specialization lattice - subclass in >1 class/subclass relationship

Top-down vs bottom up conceptual

*union

Superclass/Subclass (IS-A) relationship $S \subseteq C$

Identification - classes/objects made known by an identifier

Specialization - classifying a class of objects into more specialized subclasses

Generalization - generalizing classes into a higher class

Aggregation & Association

Ontologies & Semantic Web

Database Systems (Test 1)

Ch 1:

- * ACID (atomicity, consistency, isolation, durability)
- * DBMS has 2 functions: define database, manipulate data

Database + DBMS = Database System

Simple Complex

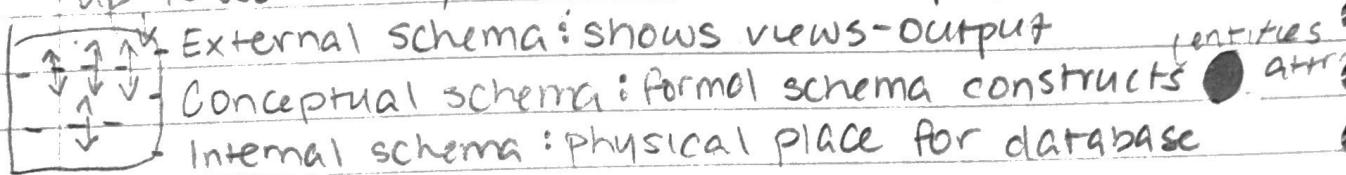
Queried	Relational	Object-relational
Not Queried	File System	Object-oriented

Ch 2:

- * High level/conceptual model: shows relationships, how the user perceives the database (ER, EER)

- * Low level/physical model: shows how data is stored, (network diagrams, paths, etc.)

- * Representational model: shows how it will be set up to work (only model with view-output) (3 schema)



- * DBMS language sets: DPL, SDL, VDL, DML

- * User interface: DB/H, GUI, MENU BASED, FORMS BASED

Ch 3:

- * ER diagrams are formal way of representing objects, properties, relationships, constraints
- Look through all symbols, frames, min/max, C, assumptions

Ch 1b:

- * Primary storage: operated by CPU, fast access, limited storage

- * Secondary storage: not processed by CPU, slower access, much larger storage capacity

- * Solid state drives (SSD): electronic circuits to store/retrieve data, store data in blocks, no mechanical components, embedded processor, read/write to set of integrated flash chips, store data when not connected to power

- * Seek time: time to position read/write head to track

- * Block transfer time: time to load block from secondary

To primary storage

- * latency: time it takes for block to rotate into position under read/write head
- * bulk transfer rate: time to transfer consecutive blocks
- * $\text{TOTAL TIME} = \text{SEEK} + \text{LATENCY} + \text{BLOCK TRANSFER TIME}$
 - large $\sim 10 \text{ msec}$
 - large $\sim 10 \text{ msec}$
 - small $\sim 1 \text{ msec}$
- * Bits \rightarrow Bytes \rightarrow Blocks \rightarrow Records \rightarrow Files \rightarrow Schema (directory)
- * voice (2,000), date/time (18 bytes)
- * variable length (unknown size - voice has 0-2,000)
- * fixed length never changes size
- * variable records happen when variable length field, multivalued attribute, field is optional, mixed file
- * block factor = $\text{floor}(\text{block size} / \text{record size})$
- * unused = block size - (block factor * record size)
- * spanned vs unspanned blocks
- * $b = \text{ceiling}(r/bf)$
- * hashing: distribute records uniformly, minimize collisions, not leave unused space
- * operations on files: retrieve/read/expand, file header carries information, searching begins at file header
- * simple selection (only rev voicemails) vs complex selection (includes several simple conditions)
- * file organization: organize data into records, blocks, ptrs
- * access methods: group of programs that allow operations on data

Ch 4 * EER model advantages: inheritance, minimize redundancy, implementation order

* open end U points to parent

* defining predicate (# of legs)

* boolean condition of predicate (# of legs = 0)

* A is disjoint, O is overlapping

* 2 reasons to make subclasses: collect special properties, participate in a special relationship

* How to determine
bottom up vs top down
what does source refer to
of assumptions

* Formal way : < PARENT > / < SUBSET >

* Unions : everything from all sets

* intersections : distinct items common to all

* superclass ; parent

* specialization - top down

* generalization - bottom up

Database Systems (9/19/2023)

Primary indexes

- * primary key, $\langle \text{value}, \text{ptr} \rangle$ pairs stored in blocks, one entry for every block, total index entries = total # of disk blocks, block anchor (first record in block), non-dense (not an index entry for every record)
- * Advantage: smaller, fewer block accesses
- * Disadvantage: insertion/deletion may cause records to be redistributed

Secondary indexes:

- * non-ordering attribute, one distinct index (candidate key), not distinct (build more than one index on same attribute), $\langle \text{value}, \text{ptr} \rangle$ pairs in blocks, dense indexes (one entry for every record), total index entries = total # of records
- * Advantage: improves search for arbitrary record
- * Disadvantage: index is dense so it requires more space and more time for searching

Clustering indexes:

- * common ordering attribute (foreign key), $\langle \text{value}, \text{ptr} \rangle$, non-dense index (one entry for every distinct value), total index entries = total distinct values
- * Advantage: search only records that hold true for "value"
- * Disadvantage: insertion/deletion are expensive

Logical indexing:

- * secondary index

Physical indexing:

- * Primary/Clustering Indexes

B-Tree

- * B for balanced, node \geq sub-node, levels (0-n)
- * Highest level = leaves
- * Internal nodes - leaves lower than leaf level
- * Level 0 is root node

Database Systems (9/28/2023)

Relational Model - words, schema constructs, symbols, mathematical statements

Terminology - relation, tuple, attribute, domain, current relational state, roles, nulls, degree

NULL - unknown but does exist, not applicable in record, doesn't exist and cannot ever be filled

* New constraints become important: no multivalue or composite attributes

* sets have no order

R - name of relation

r(A) - the relationship known as R

t_i(A_n) - value of the nth attribute in row

R.A - identifies relation with its attribute

Q, R, S - denote different relation names

grs - denotes state of entities

Q, R, S

tuv - used to denote tuples of different types

dom(R) = DOMAIN CONSTRAINTS - specify an allowable range of values

+ [S, K] ≠ [SK₂] KEYS - define relationship & constraints (elements of set are distinct)

SUPERKEY - subset of all attributes for a relation

MINIMAL SUPERKEY - smallest set of attributes

CANDIDATE KEY - deploy an index to work through records quickly

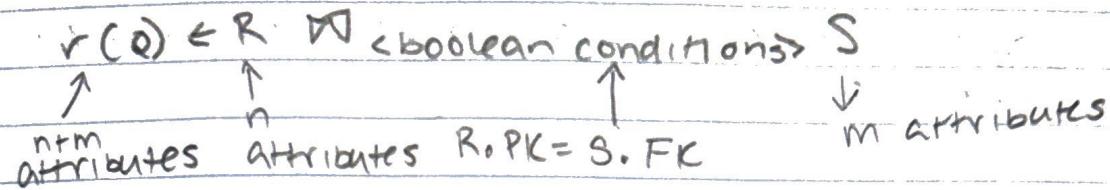
Relational database schema = set of schemas + set of integrity constraints

* Integrity of a database

↳ methods to promote better integrity: auto-filling, special key structures, displaying user entry, error messages, match, disallowing deletion of records, forcing deletion

Transaction - some executing program that includes database operations

JOIN - shows things different in relationship
* new relation returned



Theta join \bowtie any boolean equalities & inequalities $<, >, \neq, =, \leq, \geq$

Equijoin \bowtie only boolean equalities AND & OR

Natural join * only boolean equalities, duplicated values kept out AND & OR

Database Systems (10/13/2023)

left outer join

$$Q \leftarrow R \bowtie_{R_1=S_1} S$$

Returned:

- every row from R
- * matching rows from S and NULL values where S matched tuple in R

right outer join

$$Q \leftarrow R \bowtie_{R_2=S_2} S$$

Returned:

- every match from R, all other rows from S & NULL values where nothing from R matched tuple in S

full outer join

Returned:

- Every row from R with \$ without matches
- Every row from S

union

- combining two sets (no duplicates)

intersection

- rows that appear in both

difference

- keep the rows in R that aren't in S